

Human-Centric Artificial Intelligence

Implementation Report

Kartik Anand — Matriculation 676049

Hamburg University of Technology SoSe 2026

September 4, 2026

This report records how the coursework was built and, more importantly, why each open decision was settled the way it was. The course briefs leave several choices deliberately unspecified; those are the passages worth reading. Project 1 is complete and described in full. Projects 2 to 4 have their own sections, filled in as the work is done.

Contents

1	Scope, layout and environment	2
1.1	Repository	2
1.2	Environment and running it	2
1.3	Status of the four projects	3
2	Project 1 — Automated Machine Learning	3
2.1	What the brief asks	3
2.2	Structure of the app	3
2.3	Reading the data	4
2.4	Visualisation	4
2.5	Model selection	8
2.6	Models, hyperparameters and scores	10
2.7	Explaining the vocabulary in place	10
2.8	What each candidate got wrong, on request	11
2.9	Answering Task 4: who decides what	11
2.10	Results	12
2.11	Verification	14
2.12	Where the inputs are	15
3	Project 2 — Explainability	15
3.1	Tasks 1 to 3: complexity as something the user holds	15
3.2	Task 4: counterfactual explanations	18

3.3	Task 5: PDP and ALE, written from scratch	18
3.4	Keeping the regions linked	20
3.5	Verification	20
3.6	Marking what had to be our own	20
4	Project 3 — Active Learning for Learning-to-Defer	21
4.1	Outcome	21
5	Project 4 — Preference elicitation	21
5.1	Task 1: the features, and the constraint that decides them	22
5.2	Task 2: Plackett–Luce	22
5.3	Task 3: the study, and a result that reversed it	22
5.4	Task 4: the interface	23
6	Use of AI tools	23

1 Scope, layout and environment

1.1 Repository

All four projects live as separate Django apps inside one Django project, as the brief requires, and the whole thing is submitted as a git repository. The working tree sits at `~/Documents/TUHH/SoSe2026/HCAI/`:

Path	Contents
HCAI-PBL/	The Django repository. Fork of ppaamm/HCAI-PBL; remote is k-styles/HCAI-PBL.
docs/	Course overview and the project briefs.
lectures/	The twelve lecture PDFs and the demo notebooks.
data/	iris.csv and diabetes.csv, used during development.
report/	The source of this document.
.venv/	Python environment.

Inside the Django repository, `home/` is the landing page, `demos/` is the course’s example app left untouched, `pb1/` holds the settings and root URL configuration, and each project gets its own app directory.

1.2 Environment and running it

Python 3.13 with Django 6.1, scikit-learn 1.9, pandas 3.0, matplotlib 3.11; the exact versions are pinned in `requirements.txt`. From the repository root:

```
| python -m venv .venv
```

```
source .venv/bin/activate
pip install -r requirements.txt
python manage.py migrate
python manage.py runserver
```

The site is then at <http://127.0.0.1:8000/home/>.

1.3 Status of the four projects

Project	Status	Subject
1	Complete	Automated machine learning: a supervised learning interface.
2	Complete	Explainability: model complexity, counterfactuals, feature effects.
3	Complete	Active learning for learning-to-defer, with its own PDF report.
4	Complete	Preference elicitation: a movie recommender, a designed user study, and its own PDF report.

2 Project 1 — Automated Machine Learning

2.1 What the brief asks

Four tasks. Task 1: put the group members on the home page, from Python rather than from the template. Task 2: create the `project1` app and link it from the home page. Task 3: a module that uploads a CSV, reads it and visualises it. Task 4: implement the full training pipeline, and — the part that carries the marks — decide *what the user should be in control of and what should be done automatically*.

That last question is not decoration. Lecture 1 works through a ridge regression on the diabetes dataset, chooses the penalty, reports the error, and closes the slide with “Conclusion: no human needed!” — immediately followed by a slide titled “Bringing the human back in the loop”, which lists the choices that were silently made by a person at every step: who collected the data, who picked the model, who picked the evaluation criterion, who will use the result. The project is named *Automated Machine Learning* for the same reason. The application was therefore built so that the division of labour between the user and the machine is the thing you see first, not an implementation detail.

2.2 Structure of the app

File	Responsibility
<code>data.py</code>	Reading a CSV, discarding bookkeeping columns, deciding classification vs. regression.
<code>learning.py</code>	Model and score registries, the train/test split, fold construction, the sweep, the selection rule.
<code>plots.py</code>	Matplotlib figures written into <code>MEDIA_ROOT</code> .
<code>models.py</code>	<code>TrainingRun</code> : one row per sweep, so runs can be compared afterwards.
<code>forms.py</code>	Upload and training forms; the training form’s choices depend on the problem type.
<code>views.py</code>	Three pages — data, visualisation, training — plus the run log.

Roughly 1780 lines including templates and stylesheet. The uploaded file is written to `MEDIA_ROOT` under a random token and only the token is kept in the session; the rows are re-read on each request. Keeping a parsed `DataFrame` in the session would serialise the entire dataset on every request, and files of this size cost nothing to re-read.

2.3 Reading the data

The brief fixes the format — header in the first row, target in the last column — and then leaves three things open. Each was settled as an explicit rule.

Discarding an index column. The brief says datasets sometimes carry a per-row id and that one “can either assume that the provided datasets do not include such a column, or filter it out”. Matching on the column name is the obvious approach and it is wrong: `patient_id` can be a genuine feature. The rule used here is structural — integer, unique across every row, and consecutive. A column satisfying all three is numbering the rows and nothing else.

```
def looks_like_row_number(col):
    if not pd.api.types.is_integer_dtype(col):
        return False
    v = np.sort(col.to_numpy())
    if len(np.unique(v)) != len(v) or len(v) < 10:
        return False
    return bool((np.diff(v) == 1).all())
```

Classification or regression. The brief asks explicitly whether the user must state the problem type or whether the app should detect it. The answer taken here is *both*: detect it, show the verdict on screen, and let the user overrule it.

The overrule is a single click on the visualisation page, not a second upload. The file is already on disk and the problem type is only a note in the session, so requiring the same file to be uploaded again in order to change one setting would be ceremony rather than a safeguard — and it would discard the plots and the position the user had reached. Where the override is impossible the request is refused, the reason is given, and the dataset they had stays exactly as it was. A non-numeric target is labels. A numeric target is labels only if its values are whole numbers and there are few enough distinct ones to plausibly be classes. The common shortcut — a fixed cutoff of ten distinct values — misreads a twelve-class problem in one direction and a twenty-row regression table in the other, so the cutoff grows with the size of the data:

$$\text{cutoff} = \min(20, \max(10, \lfloor \sqrt{N} \rfloor)).$$

For the 150 rows of iris this gives 12; its three species fall well inside it.

Columns and rows that cannot be used. Non-numeric feature columns are dropped, as are rows with a missing target or missing features. Nothing is imputed silently. Every such action is reported on the page as a note, naming the column and the row count, because a cleaning step the user cannot see is a cleaning step the user cannot check.

2.4 Visualisation

The brief sets a floor — upload, read, visualise — and then says plainly that the more is done here, the higher the mark. The module was built to that, not to the floor. Everything is

matplotlib written into the media directory and loaded as an image, which is the route the course’s demos app demonstrates.

Two defects in that demo were not copied. Its figure filename is a fixed `myplot.png`, so two browser tabs overwrite each other’s plots; here the filename is keyed to the session token. And it never sets the `Agg` backend, which fails in a server thread; that is set before `pyplot` is imported.

What was read. Before any figure, the first ten rows are shown as the application understands them, after the dropped columns. If a file has been misparsed, every plot below it is answering the wrong question, and the user should be able to see that in one glance rather than infer it from a strange histogram.

The scatter the brief names. Two shapes are described: for classification, two chosen features coloured by class; for regression, either two features or *one feature against the target label*. Both are supported. The axes are chosen by the user, and on a regression problem the target itself appears in the axis list, at which point colour is dropped because it would only repeat an axis. Plotting a feature straight against the outcome is the plainest view of the thing being predicted, so it is the default there.

Every pair at once. A single scatter answers whether two chosen features separate the target. It cannot answer which two should have been chosen. A pair grid — every combination as a scatter, each feature’s own histogram on the diagonal — answers that in one figure (Figure 2). On a wide table this would be unreadable, so it is capped at six columns, chosen by the ranking below and named on the page.

Which features carry the answer. The features are ranked by how much each says about the target on its own. For regression this is the magnitude of the correlation. For classification it is the ratio of between-class to within-class variance — the quantity a one-way ANOVA is built on: a feature scores highly when the class means are far apart relative to the scatter inside each class.

$$F = \frac{\sum_c n_c (\bar{x}_c - \bar{x})^2 / (k - 1)}{\sum_c (n_c - 1) s_c^2 / (N - k)}$$

On iris this puts petal length and petal width an order of magnitude above the sepal measurements, which is the same conclusion the histograms reach visually. Both measures look at one feature at a time, so neither can see two features that are useless alone and decisive together; the page says so rather than letting the ranking be read as a feature selector.

All features at once. The columns are standardised and projected onto the two directions of greatest variance (Figure 3). This is written out as an SVD of the centred matrix rather than imported, since that is all principal components are. It answers a question none of the other figures can: if the classes separate here but in no single pair above, the information is spread across features rather than sitting in one, which is an argument for a model that combines them. On iris the leading component carries 73.0% of the variance and the second 22.9%, matching the published values for this dataset and serving as a check on the implementation.

The rest. Per-feature histograms split by class show where the classes overlap — on iris the petal measurements separate them and the sepal measurements do not (Figure 4). On a regression problem the target gets a panel of its own, because a skewed outcome is worth knowing about before choosing a score: squared error will chase the tail. A correlation heatmap shows which features carry the same information twice, which is exactly what the penalty term

in ridge regression exists to handle. Finally a descriptive statistics table, and for classification a class-balance table showing the proportions that the test split and every fold preserve.




Figure 1: Petal length against petal width on the iris data, coloured by species. Setosa is linearly separable from the other two on these axes; Versicolor and Virginica touch. Choosing the sepal measurements instead gives a far muddier picture — which is the point of letting the user choose the axes.




Figure 2: Every pair of features at once, with each feature's own distribution on the diagonal. Read in one pass: Setosa detaches from the other two in any panel involving a petal measurement, and Versicolor and Virginica touch everywhere. That is the whole modelling problem, before a model has been fitted.




Figure 3: All four features standardised and projected onto their two directions of greatest variance, computed as an SVD of the centred matrix. The leading component carries 73.0 % of the variance, the second 22.9 % — the published figures for standardised iris, which makes this a check on the implementation as well as a plot.




Figure 4: Per-feature distributions split by species. The two petal features separate the classes almost completely on their own; the two sepal features overlap heavily. A user who reads this before training knows what to expect from the model.

2.5 Model selection

This is the part of the project with genuine design content, and the part where the obvious library call is the wrong answer. `GridSearchCV` would do the whole of Task 4 in one line. It

would also make every decision the brief is asking about, invisibly. The pipeline below is written out instead.

Holding out a test set. The rows are split before anything else happens. For classification the split is taken class by class and every class is forced to contribute at least one test row, so a rare class cannot disappear from the evaluation entirely.

Constructing folds. The training rows are dealt into k folds by shuffling the indices of each class separately and distributing them round-robin. Every fold then carries close to the class proportions of the whole training set. This matters on small data: with iris and five folds, an unstratified shuffle can hand one fold twice as many of a species as another, and the resulting validation curve is partly measuring the shuffle rather than the model.

```
def make_folds(y, n_folds, rng, stratify):
    idx = np.arange(len(y))
    groups = [idx[y == label] for label in np.unique(y)] if stratify else [
        idx]

    folds = [[] for _ in range(n_folds)]
    for group in groups:
        for offset, position in enumerate(rng.permutation(group)):
            folds[offset % n_folds].append(position)
    return [np.array(sorted(f)) for f in folds]
```

Checked on iris with a 25% test split: 114 training rows dealt into folds of 24, 24, 24, 21 and 21, each holding exactly equal numbers of the three species, and the folds verified to be disjoint and to cover every training row.

The sweep. For each candidate value of the hyperparameter, the model is fitted k times, once per fold, and scored on the held-out fold. The mean, the standard error across folds, and — separately — the score on the rows each model was fitted on are all recorded. The last of these is not needed to pick a winner; it is recorded because the gap between the two curves is what overfitting looks like, and the user should be able to see it rather than be told about it.

Choosing a value: the one-standard-error rule. Taking the best mean is taking the maximum of k noisy numbers. The default rule here instead collects every value whose cross-validated score lies within one standard error of the best, and among those keeps the *simpler* model — the shallower tree, the larger neighbourhood, the stronger penalty. Those values are not distinguishable from the winner given the spread of the folds, so preferring the plainer one costs almost nothing on this sample and is less likely to be an artefact of one particular shuffle. There is no scikit-learn equivalent for this. The user can switch back to plain argmax from the form; both rules are offered because which one is right is a judgement, not a fact.

```
peak = (max if score.higher_is_better else min)(rows, key=lambda r: r.mean)

if score.higher_is_better:
    tied = [r for r in rows if r.mean >= peak.mean - peak.se]
else:
    tied = [r for r in rows if r.mean <= peak.mean + peak.se]

if rule == "one_se":
    order = (lambda r: r.value) if algo.simpler == "low" else (lambda r: -r.value)
    winner = min(tied, key=order)
```

```

else:
    winner = peak

```

Keeping the test score honest. The hyperparameter is chosen entirely inside the training rows. The test rows are used once, after the choice is fixed, for a single final fit and evaluation. Lecture 1’s demonstration selects the penalty on the test set and then quotes that same test error as the result; the number that comes out of that procedure has already been optimised against. The distinction is small on the diabetes data and large in principle, and it is stated on the training page so the user knows which of the two they are reading.

2.6 Models, hyperparameters and scores

One hyperparameter is swept per model, which is what the brief describes. Models whose behaviour depends on distances or on a coefficient norm are wrapped in a standardising pipeline; trees are not, because they are invariant to feature scale.

Model	Swept	Simpler end
<i>Classification</i>		
k -nearest neighbours	k	larger k — a smoother boundary
Decision tree	max depth	shallower
Logistic regression	C	smaller C — a stronger penalty
SVM, RBF kernel	C	smaller C
<i>Regression</i>		
Ridge regression	α	larger α
k -nearest neighbours	k	larger k
Decision tree	max depth	shallower

Ridge with a swept α is lecture 1’s own worked example, kept deliberately so the application can be checked against a known result. Scores are accuracy and macro F1 for classification, and R^2 , mean squared error and mean absolute error for regression. All of them are reported after every run; the one selected in the form only decides the winner. Reporting a single number would reproduce precisely the habit lecture 1 criticises — the defect-detection example, where a false negative and a false positive are not remotely the same event.

2.7 Explaining the vocabulary in place

Lecture 1 states the problem this section solves: “In practice, most users are no machine learning experts: they ignore the characteristics of the algorithms and models, and they have no understanding nor intuition on the parameters.” An interface that offers a box labelled C and no way to find out what C is has handed the user a choice they are not equipped to make, and then presented the result as though they had made it.

Thirty-three terms are written out in plain language — every model, every hyperparameter, every score, and the ideas behind them: overfitting, standard error, cross-validation, the held-out test set, the selection rule. Each has its own page. No formula appears without first being said in words, and each page ends with what to actually do.

The delivery matters more than the content. A glossary the user has to navigate to is help that arrives after the confusion has already cost them something, so the links are placed where the words are:

- **Inline, in the prose.** The sentence describing the split of responsibility reads “*Yours*: the model, the values to try, how much data is held back, how many folds, the score, and how ties are broken” — and each of those six phrases is itself a link, marked with a dotted underline. Hovering shows the one-line answer without leaving the page at all. There are 36 such links on the training page and 20 on the visualisation page.
- **Beside form labels.** A small outlined marker next to each field. Three of them are rewritten by script as the dropdowns change: select a decision tree and the marker beside “values to try” points at the page on tree depth; select k -nearest neighbours and it points at the page on k ; change the score and its marker follows.

2.8 What each candidate got wrong, on request

The confusion matrix for the winning model answers “where was it wrong”. It cannot answer the more useful question, which is how the mistakes *changed* as the hyperparameter moved.

Every row of the sweep table can therefore be expanded to show its own error breakdown. The predictions come from the cross-validation that has already happened: each fold’s model predicts the fold held out from it, and those out-of-fold predictions are collected per candidate. They cost nothing extra to compute, and no candidate is ever scored on rows it was fitted on — so this is available for every value without spending the test set on any of them. For regression, where no confusion matrix exists, the same panel reports the typical miss, the worst miss, and whether the model leans high or low.

The panels are collapsed by default. Nine extra tables opened on arrival would bury the result; the question only arises once the user has a reason to ask it.

The iris sweep shows why it is worth asking. A tree of depth 1 scores 0.6667, which merely says “poor”. Expanding it says what poor means here: Setosa 38 of 38, Versicolor 38 of 38, Virginica 0 of 38. The one-question tree never predicts the third species at all. That is a different and far more actionable fact than a single number, and it is invisible in the score column.

2.9 Answering Task 4: who decides what

The user decides	The app does
The model	Builds class-balanced folds
The values to sweep	Fits every candidate
The size of the test set	Keeps the test rows out of the selection
The number of folds	Reports every applicable score, not just one
The score to select on	Skips values that cannot be fitted, and says which
How ties are broken	Records the run in the log
The random seed	

The split is deliberate. Everything on the left is a judgement that depends on what the user is trying to achieve; everything on the right is mechanical, and automating it removes drudgery rather than authority.

Next to the form sits a second button, *Decide for me*, which runs every model for the detected problem type over its default grid and keeps the winner, asking nothing at all. It then shows what it decided — including the models it discarded and their scores — with a note saying that a model further down the table may be easier to explain or cheaper to run, for a difference in score too small to be worth the trade. On the iris data the automatic route selects k -nearest neighbours with $k = 13$ and reaches 0.9167 test accuracy, while a decision tree of depth 3, chosen manually, reaches 0.9444 and can be read off in three questions. That contrast is the argument of lecture 1 rendered as two buttons on one page.

Every run — manual or automatic — is written to a `TrainingRun` row and listed at `/project1/history/`, so two sets of settings can be compared after the fact instead of from memory. This is the use of Django models the brief suggests for projects that offer several learning algorithms.

2.10 Results

Iris, decision tree, accuracy. Table 1 is the case the selection rule exists for. Nine depths were swept; depths 3 through 14 are *identical* in cross-validation while the training score climbs to a perfect 1.000. Plain argmax has nothing to separate them. The one-standard-error rule keeps depth 3, which reaches 0.9444 accuracy and 0.9441 macro F1 on the 36 held-out rows, and which a person can actually read.

Max depth	CV accuracy	$\pm$ SE	On training rows	
1	0.6667	0.0000	0.6667	
2	0.9476	0.0170	0.9759	
3	0.9655	0.0158	0.9847	kept
4	0.9655	0.0158	0.9956	tied
5	0.9655	0.0158	0.9978	tied
6	0.9655	0.0158	1.0000	tied
8	0.9655	0.0158	1.0000	tied
10	0.9655	0.0158	1.0000	tied
14	0.9655	0.0158	1.0000	tied

Table 1: Five-fold sweep over tree depth on iris, 114 training rows. Everything from depth 3 onwards is the same model as far as cross-validation can tell; the only thing that changes is how much of the training set it has memorised.




Figure 5: The same sweep as a curve. The solid line is cross-validation with its one-standard-error band; the dashed line is the score on the rows each model was fitted on. Where the two separate, the model is memorising rather than learning.

Diabetes, ridge regression, mean squared error. Here the rule actually overrules the peak. The best mean sits at $\alpha = 1$, but $\alpha = 100$ lies within one standard error of it, and the far more strongly penalised model is kept. On the held-out rows it scores marginally *better* — 2889.2 against 2893.8 — which is luck rather than proof, but it is not the trade going wrong.

α	CV MSE	$\pm$ SE	On training rows	
0.001	3026.80	183.48	2840.24	tied
0.01	3026.73	183.56	2840.24	tied
0.1	3026.12	184.26	2840.29	tied
1	3024.18	188.70	2842.79	best mean
10	3024.72	194.38	2861.22	tied
100	3082.98	198.44	2981.10	kept
1000	4150.11	167.53	4107.01	

Table 2: Five-fold sweep over the ridge penalty on the diabetes data. The final test MSE of 2889.2 sits alongside lecture 1’s reported 2856 for the same dataset and model — close enough to treat the pipeline as validated against a known result, and different enough to be explained by the split and by the fact that the lecture selects its penalty on the test set.




Figure 6: Ridge on the diabetes data, scored by R^2 . The curve is almost flat across four orders of magnitude of α before it collapses — which is the honest shape of this problem, and is invisible if a grid search returns only its argmax.

2.11 Verification

A script exercises every view through Django’s test client: both datasets, all seven model and rule combinations, custom grids, the automatic route, the run log, and ten error paths. Thirty-four checks, no failures, no server errors and no browser console errors. The generated figures were inspected directly rather than assumed, and the two-dimensional projection was checked against the published variance shares for standardised iris.

Four defects were found this way and fixed:

1. A three-row file reported “no usable numeric features” instead of “too few rows”, because the index-column filter consumed its only feature column first. The row-count guard now runs before any column is discarded, and the filter requires at least ten rows before it will call a column bookkeeping.
2. On a small dataset the default k -nearest-neighbours grid contained values of k larger than the number of rows a fold could train on, and scikit-learn raised. Candidates that cannot be fitted are now filtered out before the sweep and named on the page, rather than aborting the run.
3. The warning shown after an automatic run asserted that “a decision tree of depth 3 would have been almost as good” — true for iris and false in general. It was replaced with text that points at the leaderboard instead of hardcoding a result.
4. Forcing the problem type to regression on a dataset whose target is text returned a server error, because nothing checked that the override was possible before the feature ranking tried to correlate against the string `Setosa`. The override is refused at load time now, with a message naming the column and what to do instead.

The fourth is the instructive one. The override had been tested before release — but only on a numeric target, which is the one case where forcing regression happens to work. A test that

exercises a feature on the single input it was written against confirms nothing. The problem type is now checked across every combination of the three settings and three shapes of target column: text, continuous, and small whole numbers where either reading is defensible. The opposite direction, forcing a continuous target into classes, was found to be technically caught already but to report itself by listing two hundred singleton “classes”; it now says that the column reads as a quantity and suggests regression.

2.12 Where the inputs are

Location	What to change
home/views.py, GROUP	Names and matriculation numbers. Set in Python, as Task 1 requires.
The upload form at /project1/	The dataset, and the problem type if the detected one is wrong.
project1/learning.py, ALGORITHMS	Default hyperparameter grids, if the form’s per-run field is not enough.

3 Project 2 — Explainability

Palmer Penguins: 344 rows, of which 333 are complete. Four numeric measurements (bill length, bill depth, flipper length, body mass) and three categorical columns (island, sex, year) predict the species. That split between numeric and categorical is what Tasks 4 and 5 turn on.

One stratified split of 233 training and 100 test rows is made once and shared by every model on the page. Task 2 compares models by test accuracy, so a model must not be able to look better merely by having been handed an easier test set.

3.1 Tasks 1 to 3: complexity as something the user holds

Equation (1) of the brief trades empirical loss against a complexity measure Ω . Task 2 makes that trade interactive: a slider on λ selects whichever trained model maximises $\text{acc}_{\text{test}} - \lambda\Omega(f)$.

Reading the tree back out. Task 1 asks for the tree to be shown. `plot_tree` is the conventional answer and is included, but it is also the first thing to stop being readable — past a dozen leaves it is a wall of small boxes. Since the subject of this project is precisely whether a person can follow the model, the fitted tree is also walked directly and rendered as one sentence per root-to-leaf path. At $\lambda = 0.02$ the whole model is three lines:

Rule	Support
flipper ≤ 207.5 and bill length $\leq 43.35 \rightarrow$ Adelie	100 penguins, 98 % pure
flipper $> 207.5 \rightarrow$ Gentoo	87 penguins, 94 % pure
flipper ≤ 207.5 and bill length $> 43.35 \rightarrow$ Chinstrap	46 penguins, 93 % pure

One detail matters for readability. The model is fitted on one-hot columns, so a split reads internally as `island=Biscoe  $\leq 0.5$` . Printing that would be faithful to the arithmetic and useless to a reader, so categorical splits are put back into words as “island is not Biscoe”.

Choosing Ω for logistic regression (Task 3). The brief fixes Ω for a tree as the leaf count and leaves the linear case open. The measure used here is **the number of original features carrying a non-zero coefficient**, under an L1 penalty.

The reasoning is that Ω should measure the same thing in both cases: how much of the model a person has to read before they can say what it does. For a tree that is the number of questions it may ask. For a linear model it is the number of features that must be looked up, and a coefficient of zero costs the reader nothing. L1 is what makes the count move at all — an L2 penalty shrinks coefficients without ever reaching zero, which would leave Ω pinned at the full feature count for every setting and make the slider inert. Original features are counted rather than encoded columns: “island” is one thing to go and find out, whether it expands into two dummy columns or ten. Task 4 referring to the slider as “sparsity” points the same way.

Which models the slider can actually reach. For a fixed candidate, $\text{acc} - \lambda\Omega$ is a straight line in λ with slope $-\Omega$. Maximising over candidates is therefore taking the upper envelope of a family of lines, and only candidates whose line touches that envelope are ever selected. Plotted as (Ω, acc) , that envelope is the upper concave hull: anything strictly below it loses at every λ to something simpler, something more accurate, or a blend of the two.

This is worth computing rather than leaving implicit, because it converts the slider from a dial to be dragged blindly into a short list with exact boundaries. Of the ten distinct trees trained, *four* can ever be selected; the 5-, 6-, 7-, 8-, 10- and 13-leaf trees are unreachable at every λ even though they look perfectly reasonable in the results table.

λ from	up to	Ω	accuracy	
0.000	0.005	12	0.980	most accurate available
0.005	0.010	4	0.940	
0.010	0.140	3	0.930	the three-rule tree above
0.140	—	2	0.790	one question, two answers

Table 3: Every model the tree slider can produce, and the exact range of λ that produces it. The widest band by far is the three-leaf tree, which is also the one a person can read.




Figure 7: Accuracy against complexity for the ten trees. Filled points are on the upper concave hull and are reachable; grey points are dominated and no setting of the slider produces them.




Figure 8: The same thing as lines. Each model is $\text{acc} - \lambda \Omega$; the slider picks whichever is highest at the chosen λ . The greyed lines spend their whole lives underneath someone else’s, which is what “unreachable” means.

The logistic family behaves differently and instructively: every one of its five distinct models is on the hull, so the slider genuinely sweeps from all seven features at 0.99 accuracy down to a model using none at all and scoring 0.44 — which is the majority-class rate, exactly as it should be.

3.2 Task 4: counterfactual explanations

The brief’s recipe is followed directly — sample around x , keep the draws the model assigns to the desired class, rank by MAD-weighted L^1 distance, show the closest few — with three decisions that the brief hands over.

Noising categorical data. The brief asks the question explicitly. Numeric features get Gaussian noise with standard deviation $\text{spread} \times \text{MAD}(\text{feature})$, so one spread setting means the same thing to body mass in grams and bill depth in millimetres, and it is the same scale the ranking distance uses. Categorical features cannot be nudged — there is nothing between Biscoe and Dream — so each is left alone with probability 1 – flip and otherwise redrawn from the values actually observed. Drawing from the empirical distribution rather than uniformly keeps rare combinations rare, so a proposal no penguin would ever be does not get ranked highly.

Distance across mixed types. A categorical change counts 1 and no change counts 0, which puts it on the same order as a typical numeric change once the numerics are divided by their own MAD.

Sparsity, which turned out to matter most. Perturbing all four measurements on every draw means every counterfactual returned has all four changed, and “alter everything slightly” is close to useless as advice. Each draw therefore first picks how many numeric features it may touch — uniformly from one to all four — and leaves the rest exactly as they were. Sparse proposals then exist to be found, and since changing fewer features costs less distance they win the ranking whenever they work at all.

The effect is large. Before this change, the closest counterfactual for penguin #0 towards Chinstrap altered all four measurements at distance 1.79. After it, the answer is a single change — bill length from 39.1 to 43.6 mm, distance 0.96 — which is an explanation someone can act on.

If nothing is found the search widens over four rounds, growing the number of draws and the noise. Reaching Gentoo from an Adelie needs the second round and returns a longer answer: grow the flippers past 207 mm *and* be on Biscoe. Both are biologically real distinctions in this dataset, which is a reassuring sign that the search is finding structure rather than noise.

3.3 Task 5: PDP and ALE, written from scratch

The brief requires these to be implemented rather than imported. They are computed from the fitted model’s predicted probabilities and numpy arithmetic: there is no `sklearn.inspection`, no `partial_dependence`, and no ALE package anywhere in the file.

PDP. For each grid value v , every row in the data has the feature forced to v , the model is asked, and the answers are averaged:

$$\text{PDP}_c(v) = \frac{1}{n} \sum_{i=1}^n P(c | x_i \text{ with } x_{ij} := v).$$

Grid points are placed by quantile so they follow the data rather than the axis.

ALE. Following the lecture’s definition, with the integral estimated per bin and the curve centred to zero mean over the data. The expectation is conditional — inside each bin only rows whose own value falls in that bin contribute — which is the entire advantage over the PDP.

On this dataset flipper length and body mass move together, so forcing a flipper to 230 mm while leaving body mass alone asks the model about an animal that does not exist. The PDP averages over those; the ALE never constructs them.

Exact or discretised (the brief’s question). *Logistic regression admits exact derivatives.* The scores are linear in the standardised inputs, so $\partial s_c / \partial x_j = W_{cj} / \sigma_j$, and pushing that through the softmax gives

$$\frac{\partial p_c}{\partial x_j} = p_c \left(\frac{W_{cj}}{\sigma_j} - \sum_{c'} p_{c'} \frac{W_{c'j}}{\sigma_j} \right),$$

which is written down in closed form and integrated across each bin. A class gains probability only in so far as its own score rises faster than the probability-weighted average, which is why one species’ ALE curve must fall when another’s rises.

A decision tree does not. It is piecewise constant: its derivative is zero almost everywhere and undefined exactly on the split points, so there is nothing to evaluate pointwise and finite differences across bin edges are forced. Both paths are implemented and the interface states which one produced the curve on screen.

That the two agree is the check that makes the derivative trustworthy. Computing both ways for the logistic model on all four features, the curves match to 1.5×10^{-4} at worst — a wrong derivative would diverge immediately.




Figure 9: Bill length under the three-leaf tree. Both curves step at 43.35 mm, which is exactly the tree’s split point — the correct shape for a piecewise constant model, and a useful sanity check on the implementation. Gentoo is flat because this tree separates Gentoo on flipper length instead.




Figure 10: The same feature under logistic regression, where the ALE integrand is the closed-form softmax derivative. The transition is smooth rather than stepped, as a differentiable model requires.

3.4 Keeping the regions linked

The brief requires the counterfactual and feature-effect regions to be driven by the model chosen in Tasks 1 to 3. The model class and λ live in the session and are set by one control strip that appears, identically, at the top of all three regions. Switching from tree to logistic regression changes which models the frontier offers, which model generates counterfactuals, and which route the ALE takes — all three were checked to move together.

3.5 Verification

A script drives every region through Django’s test client: both model classes at four values of λ , all three species as counterfactual targets from four different penguins, all four features under both models, and the linkage between regions. Every case passes. Alongside it:

- The hull was checked against brute force. Sweeping λ over 35,000 values and taking the arg max directly reproduces exactly the set of models the hull predicts, and the λ bands agree with the arg max everywhere except at a single floating-point tie on a boundary.
- The PDP was checked against an independent hand computation at a chosen grid point.
- The two ALE routes were checked against each other, as described above.
- Fixing the solver seed was necessary rather than cosmetic: `saga` is stochastic, and without `random_state` the fitted coefficients — and so Ω , and so the whole frontier — differed between runs. This was caught by running the same computation in three separate processes and comparing.

3.6 Marking what had to be our own

The briefs single out particular components as things to implement or decide ourselves. Rather than leave that as a claim, the requirements are held in one registry (`home/own_work.py`) which records the brief’s exact words, the file that satisfies each one, and what was done. Each such block in the code carries a marker comment, and the page at `/home/own-work/` re-reads the files and reports whether every marker is still present — so deleting or moving a marked block shows

up as a failure rather than passing quietly. The relevant entries are also shown in place, at the top of the region they govern.

4 Project 3 — Active Learning for Learning-to-Defer

Complete, with its own 9-page PDF report at [project3/report/project3-report.pdf](#), downloadable from the project page as the brief requires. The summary below states the brief; the report carries the results and the argument.

The dataset is AG News: label news articles by topic. The system is allowed to answer or to defer to a human expert, and the point of the project is to learn *when* deferring is worth it — without being given the expert’s labels up front. Four required tasks and one optional one:

1. Train a classifier on the fully labelled training set and report test accuracy. This is the baseline the human-AI team has to beat.
2. Design and implement at least one simulated expert. It must be imperfect and, specifically, competent in particular regions of the input space rather than uniformly noisy — otherwise there is nothing for a deferral policy to learn.
3. Implement a learning-to-defer strategy, evaluated on the quality of the deferral decisions and not only on accuracy.
4. Choose an active learning strategy for deciding which examples to show the expert, given no expert labels at training time. Justify the choice.
5. (Optional) An interface where a person plays the expert.

4.1 Outcome

A classifier right 92.3 % of the time is paired with a simulated expert right 69.9 % of the time; deciding per article which should answer raises the pair to 94.1 % while consulting the expert on 16 % of articles.

Two decisions carried the project. The simulated expert’s competence is a function of the input *text*, not of the true label — an expert reliable wherever the label happens to be Sports could never be discovered by Task 4, because at query time the label is exactly what is unknown. And every downstream target uses out-of-fold predictions: the classifier makes 3,403 errors in-sample against 11,350 out-of-fold, so the in-sample figure would have understated failure by $3.3\times$ and taught the deferral model that the classifier almost never needs help.

A null result was found and kept. The second simulated expert shows no separation between any query strategy, which narrows the claim from “choosing questions beats random” to “choosing questions helps when the expert’s competence is visible in the available features”.

5 Project 4 — Preference elicitation

Complete, with its own 12-page PDF report at [project4/report/project4-report.pdf](#) covering Tasks 1 to 3, downloadable from the project landing page alongside the entry point to the study itself — the two-door arrangement the brief specifies.

The subject is a movie recommender that adapts quickly to a new user, built on the IMDB 5000 Movie Dataset — metadata for roughly 5000 films, with no user ratings. A user’s preferences are modelled by a utility $U(x) = w^\top x$, where x is a feature vector for a film and w is the user’s latent preference vector to be estimated from a small number of interactions.

5.1 Task 1: the features, and the constraint that decides them

The representation is 20-dimensional over a pool of 2,734 films (those with at least 25,000 IMDB votes, since a participant comparing two films they have never heard of is guessing from metadata rather than expressing a preference).

The argument is that the feature set is governed by the *elicitation budget*, not by descriptive power. With w fitted from roughly twenty-five responses, a one-hot over the 2,398 directors is unidentifiable — two random films share a director well under once in a thousand times, so essentially every response multiplies those weights by zero. A second rule excludes anything the participant cannot see on the card: a feature that was invisible cannot have caused the choice, but the fit will credit it anyway.

A genre earns a dimension only when two random films actually differ on it, which happens with probability $2p(1 - p)$; requiring one informative pair in ten puts the cut at $p \geq 0.0528$ and admits fourteen genres. The threshold is computed in `project4/data.py` rather than hardcoded, so the prose and the code cannot drift apart.

5.2 Task 2: Plackett–Luce

A ranking is read as a sequence of choices — favourite of ten, then of the remaining nine — each stage being Bradley–Terry widened past two alternatives. The product of the stages is Plackett–Luce.

Three justifications, of which the middle one is the load-bearing one. It reduces to Bradley–Terry exactly at $n = 2$. Its pairwise marginals *are* the Bradley–Terry probabilities, so both study arms estimate the same w and can legitimately be scored against a common held-out set — without which the comparison would be a category error. And the tempting alternative of splitting a ranking into all 45 implied pairs is rejected on two grounds: it counts dependent comparisons repeatedly, and it is not a distribution over the $n!$ orderings, so it is not a likelihood at all.

5.3 Task 3: the study, and a result that reversed it

A 150-participant, between-subjects, time-matched, pre-registered design. Both arms get the same eight-minute budget rather than the same number of tasks, since what a recommender spends is the user’s patience. Because w is latent and no distance to it is computable, the primary outcome is held-out prediction: an identical block of ten fresh pairwise comparisons, scored by log loss.

The project began with the hypothesis that ranking wins, which is what the per-task Fisher information says — one ranking carries about $60\times$ one comparison. The power simulation written to size the study overturned it. Fixing time rather than task count means eight minutes buys 80 comparisons or 5 rankings, and the accumulated traces then come out at 118 against 129, a ratio of 1.09. And the trace is the wrong summary: over a session the log-determinant is 17.6 nats lower for ranking and its smallest eigenvalue $6.4\times$ smaller, because 80 comparisons touch 160 films while 5 rankings touch 50, and the 45 comparisons inside one ranking all lie in

the span of that set's ten feature vectors. Ranking piles up information about a few directions and leaves the rest to the prior.

Across 120 simulated participants per condition the arms tie when ranking is assumed flawless ($d = 0.00$) and pairwise pulls clear as soon as later positions get careless ($d = -0.34$ and -0.51). H1 was rewritten in the opposite direction, which is what pre-registration exists to make visible.

5.4 Task 4: the interface

Consent and a stratification question, then blocked random assignment, then the assigned design under a visible timer, then the held-out block, then a debrief showing the participant their own fitted taste profile and five recommendations.

Responses go to database rows rather than the session, because session-only storage loses everything when a tab closes. A pairwise choice is stored as a ranking of length two, so the schema has one representation and the fitting code never branches on the arm — the practical payoff of Task 2's first justification. Every draw derives from one stored seed per participant, so the held-out block is fixed for them and cannot be rerolled by reloading.

Ranking supports drag-and-drop *and* per-row arrow buttons, and the form serves its hidden inputs in the presented order so that scripting-disabled participants submit something coherent. This is not politeness: a drag-only interface excludes participants invisibly, since they appear in the data as abandonments and inflate exactly the secondary statistic being measured.

Twenty-four plain-language explanation pages are linked inline at the point each term first appears, participant-facing screens included — consent that requires a statistics degree to understand is not informed consent.

6 Use of AI tools

The implementation was drafted with a large language model and then reviewed. The repository contains `AI_USAGE.md`, which records component by component what that applies to. Both that file and this report are readable from inside the application itself at `/project1/notes/` — a convenience during development, gated on `DEBUG` so the page and the file it serves both return 404 in any other setting, and nothing has to be remembered and removed before submission. In short: the Django wiring, the templates, the stylesheet and the matplotlib figure code were drafted by the model; the three decisions in Project 1 with more than one defensible answer — the index-column rule, the problem-type cutoff, and the fold construction with its selection rule — were settled against the brief and lecture 1 first, and the model was used to write them down rather than to choose them. That distinction is the reason the project does not simply call `GridSearchCV`, which is what a model asked to “pick a hyperparameter” will reach for, and which is the answer the brief is designed to get past.